

Deep Learning & Applied AI

Geometric deep learning

Emanuele Rodolà
rodola@di.uniroma1.it



SAPIENZA
UNIVERSITÀ DI ROMA



Audio signals



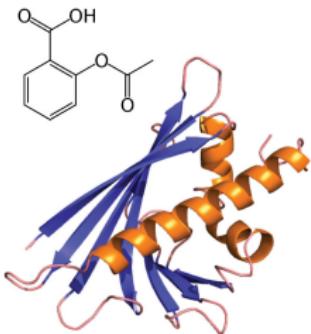
Images



Audio signals



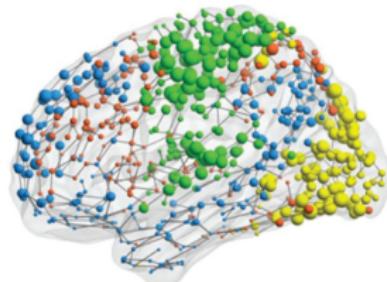
Social networks



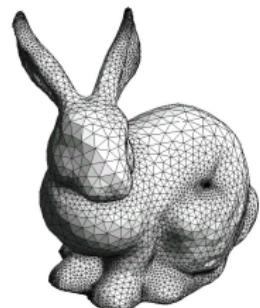
Molecules



Images

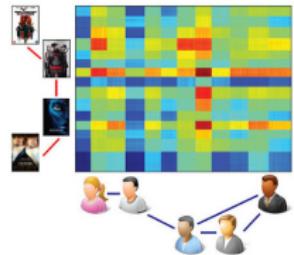


Functional networks

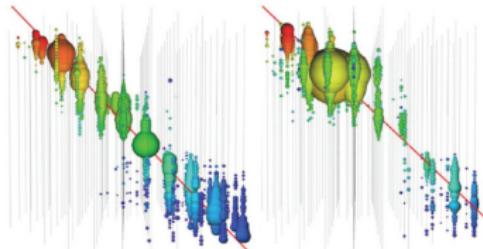


3D shapes

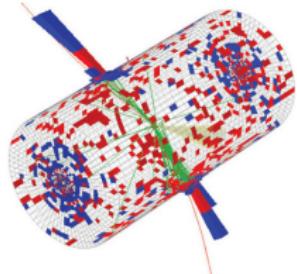
Applications of geometric deep learning



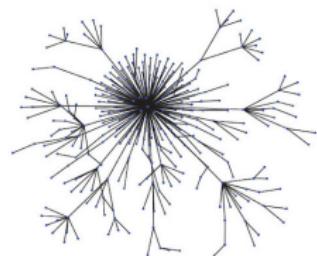
Recommender system



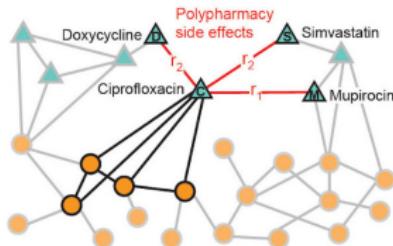
Neutrino detection



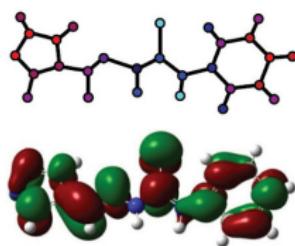
LHC



Fake news detection



Drug repurposing



Chemistry

Prototypical non-Euclidean objects



Manifolds

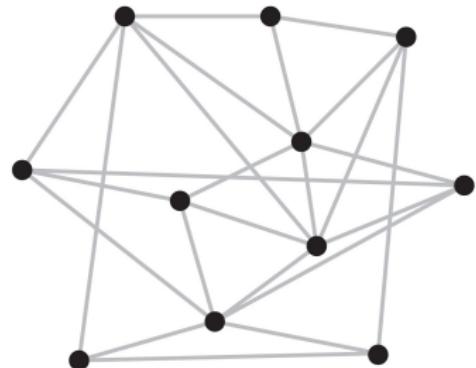


Graphs

Domain structure vs Data on a domain

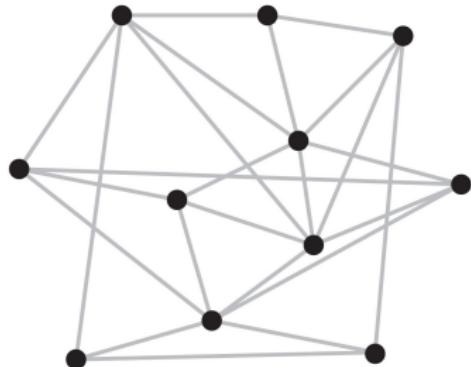


Domain structure vs Data on a domain

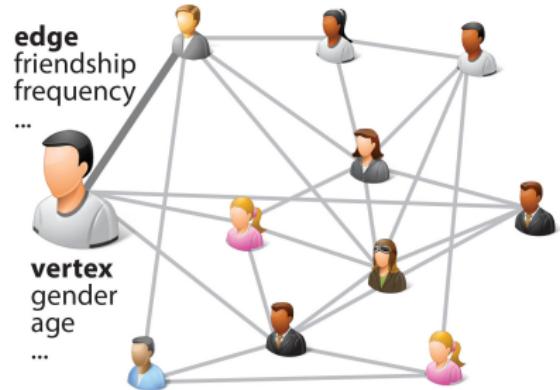


Domain structure

Domain structure vs Data on a domain

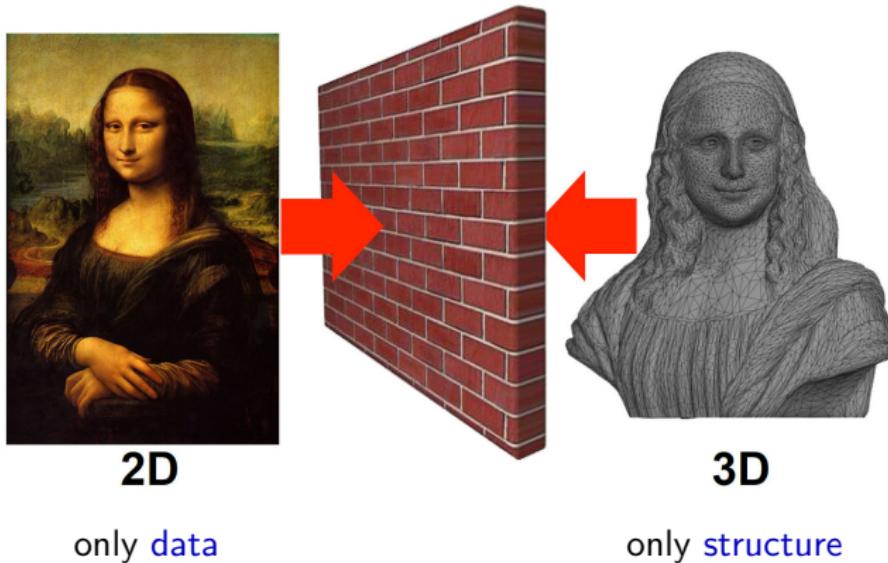


Domain structure



Data on a domain

Domain structure vs Data on a domain



Fixed vs different domain

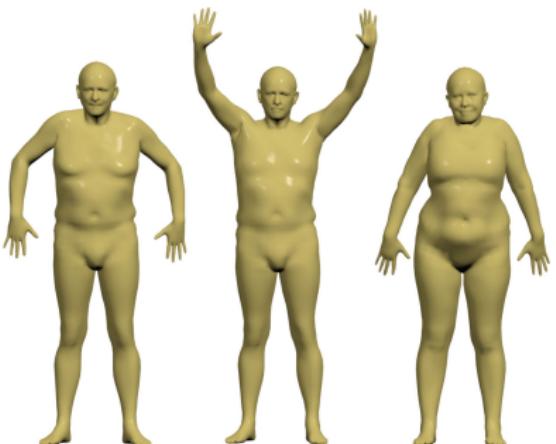


Social network
(fixed graph)

Fixed vs different domain



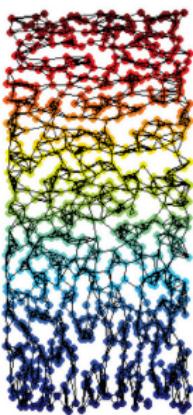
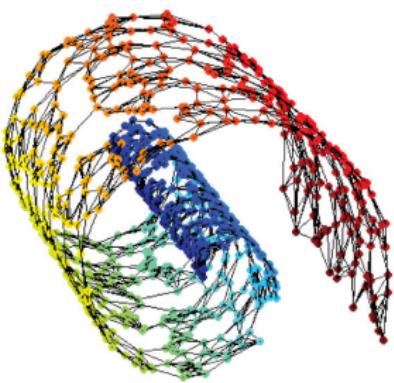
Social network
(fixed graph)



3D shapes
(different manifolds)

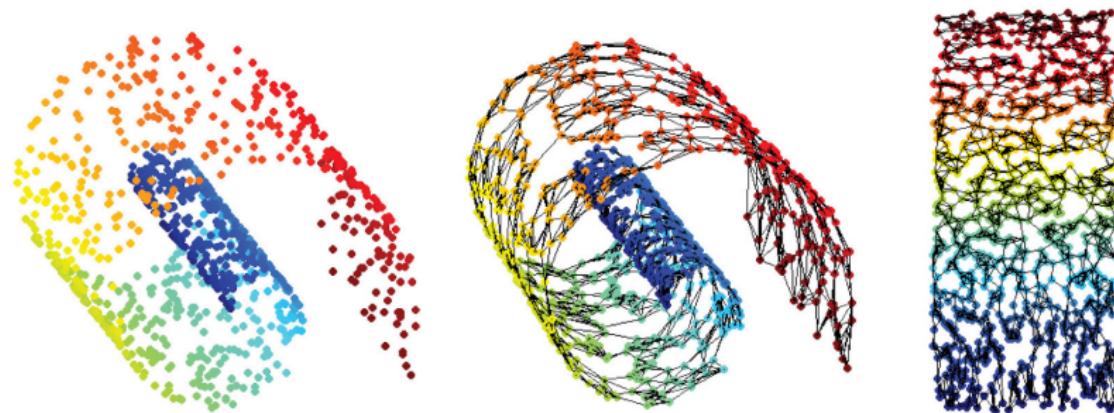
Geometric learning \neq Manifold learning

In **manifold learning**, we seek for a (possibly high-dimensional) manifold that justifies a given set of data points:



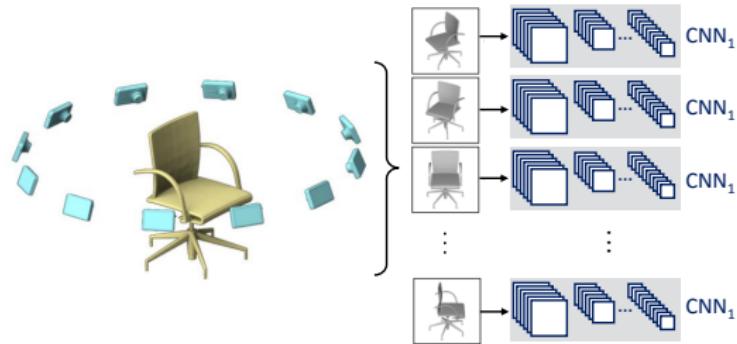
Geometric learning \neq Manifold learning

In **manifold learning**, we seek for a (possibly high-dimensional) manifold that justifies a given set of data points:



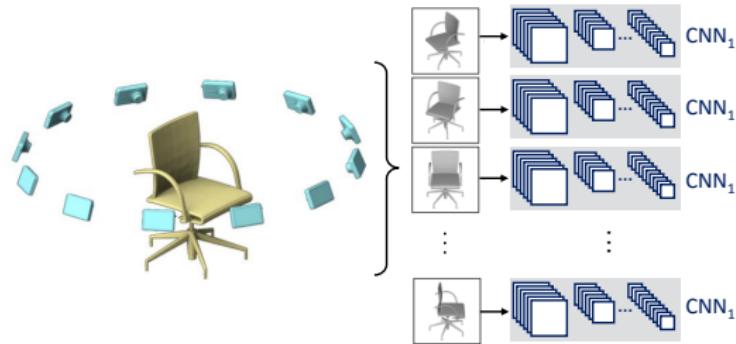
In geometric deep learning, the data has a known geometric structure.

Multi-view CNNs



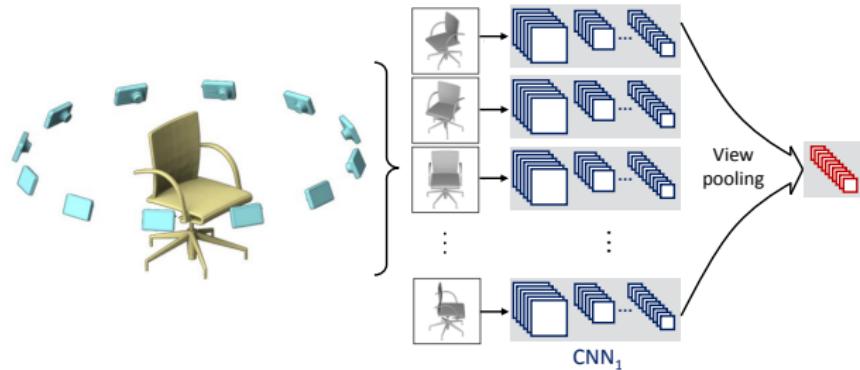
- Represent 3D object as a collection of range images

Multi-view CNNs



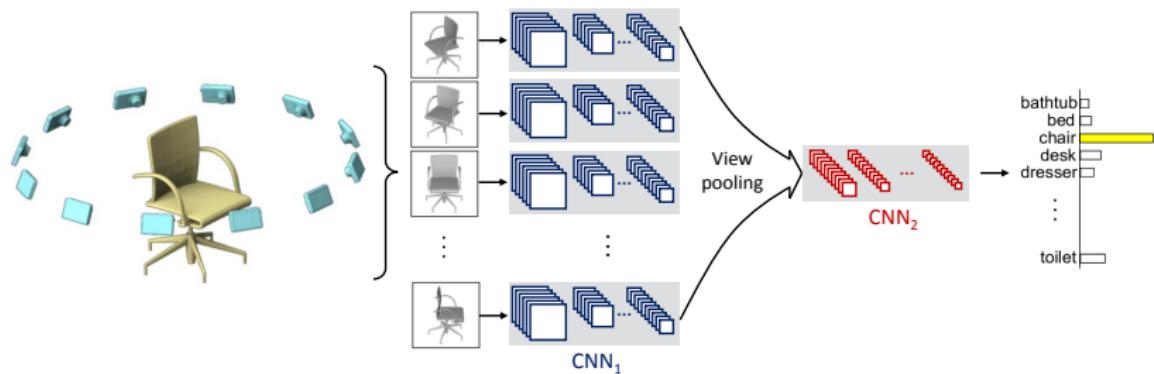
- Represent 3D object as a collection of range images
- CNN₁: Extract image features (parameters are shared across views)

Multi-view CNNs



- Represent 3D object as a collection of range images
- CNN_1 : Extract image features (parameters are shared across views)
- Element-wise max pooling across all views

Multi-view CNNs



- Represent 3D object as a collection of range images
- CNN_1 : Extract image features (parameters are shared across views)
- Element-wise max pooling across all views
- CNN_2 : Produce shape descriptors + final prediction

Applications of Multi-view CNNs

- 3D shape classification and retrieval
 - Pre-trained on ImageNet
 - Fine-tuned on 2D views



Applications of Multi-view CNNs

- 3D shape classification and retrieval
 - Pre-trained on ImageNet
 - Fine-tuned on 2D views
- Sketch classification
 - Mimic views by jittering



Applications of Multi-view CNNs

- 3D shape classification and retrieval

- Pre-trained on ImageNet
- Fine-tuned on 2D views



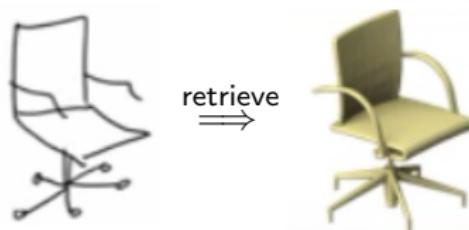
- Sketch classification

- Mimic views by jittering



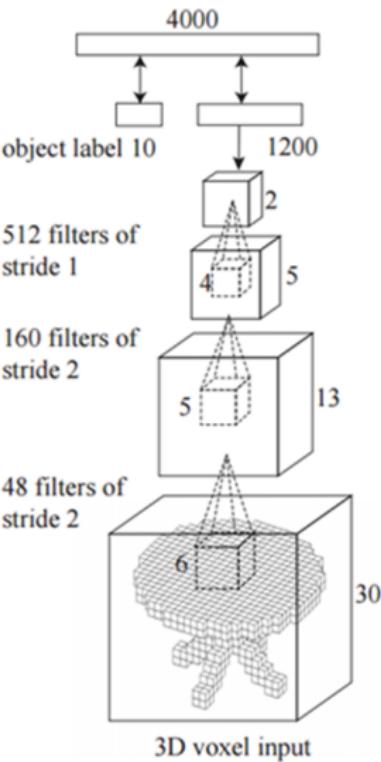
- Sketch-based shape retrieval

- Render views with hand-drawn style (edge maps)



3D ShapeNets

- Volumetric representation (shape = binary voxels on 3D grid)

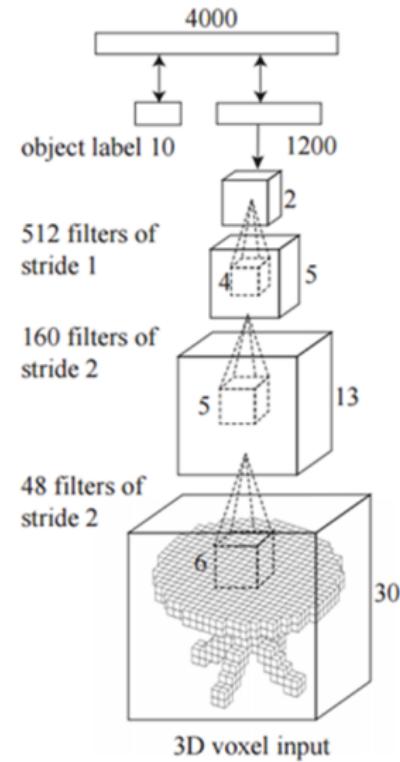
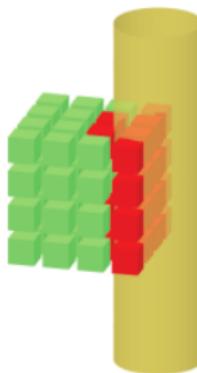


Convolutional deep belief network

Wu et al, "3D ShapeNets: A Deep Representation for Volumetric Shapes" 2015

3D ShapeNets

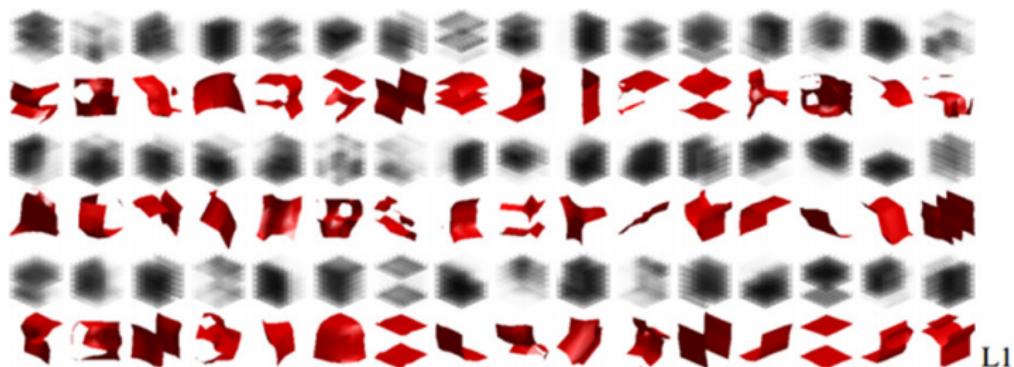
- Volumetric representation (shape = binary voxels on 3D grid)
- 3D convolutional network



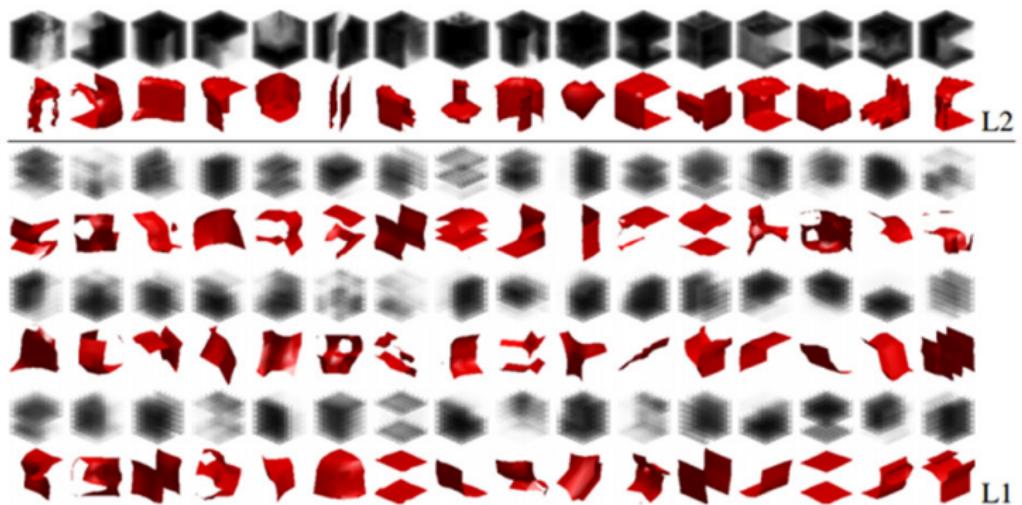
Convolutional deep belief network

Wu et al, "3D ShapeNets: A Deep Representation for Volumetric Shapes" 2015

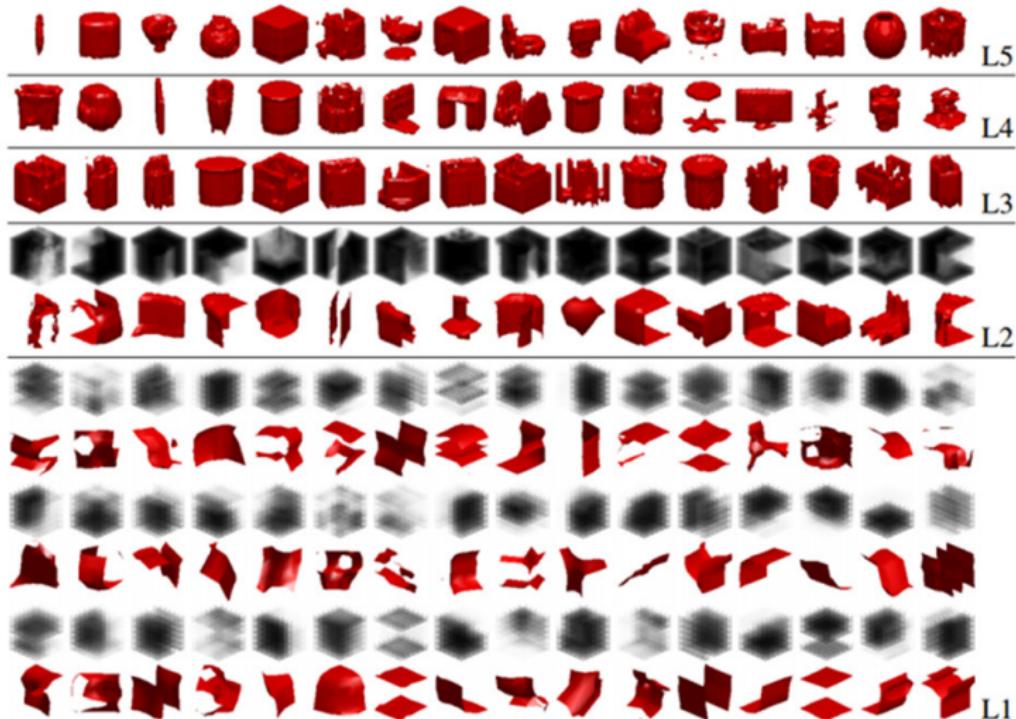
Learned features: 3D primitives



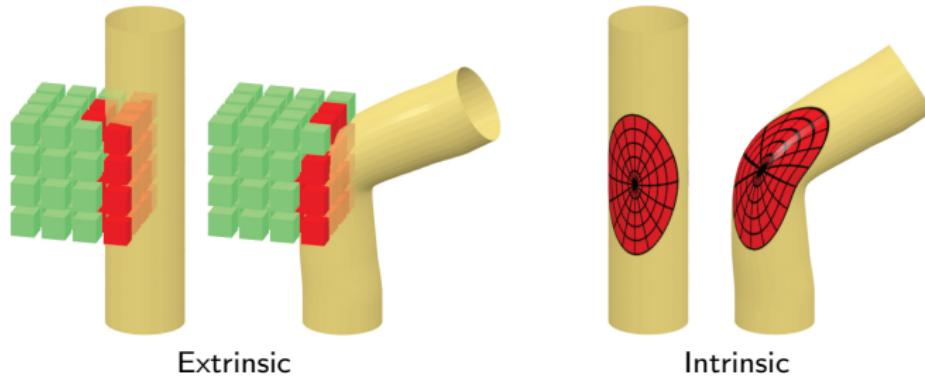
Learned features: 3D primitives



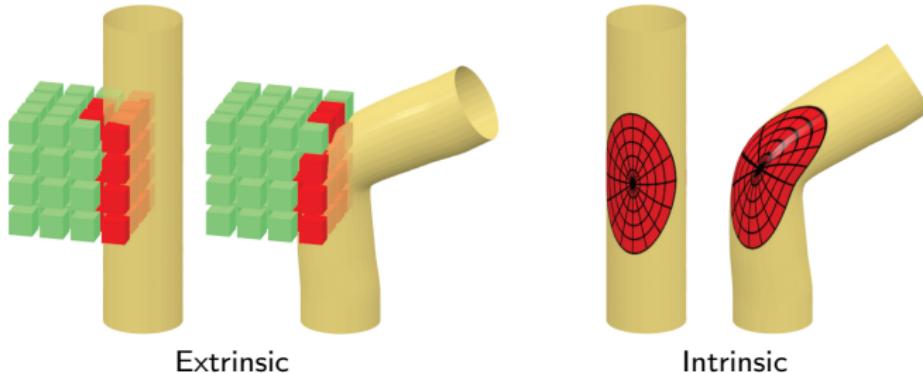
Learned features: 3D primitives



Challenges of geometric deep learning



Challenges of geometric deep learning



- How to define convolution?
- How to do pooling?
- How to work fast?

Extrinsic vs Intrinsic

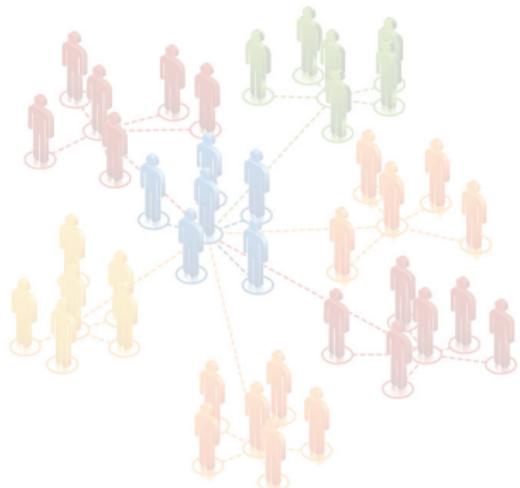
Extrinsic

Intrinsic

Prototypical non-Euclidean objects

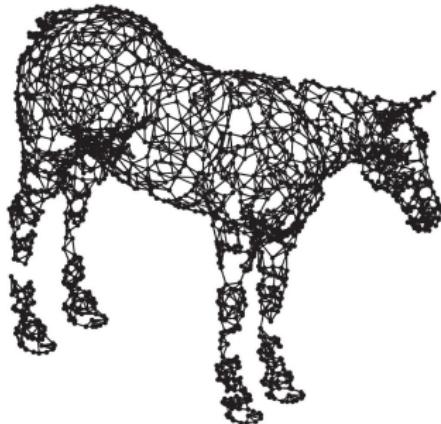


Manifolds

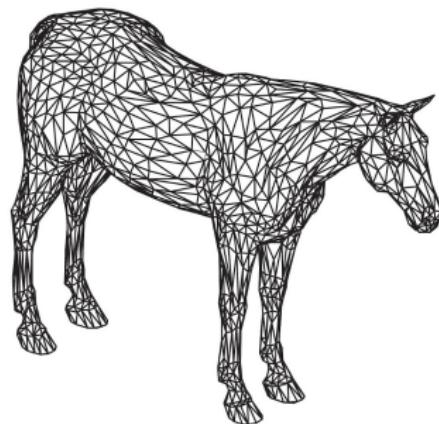


Graphs

Discrete manifolds



Nearest neighbor graph



Triangular mesh

Vertices $\mathcal{V} = \{1, \dots, n\}$

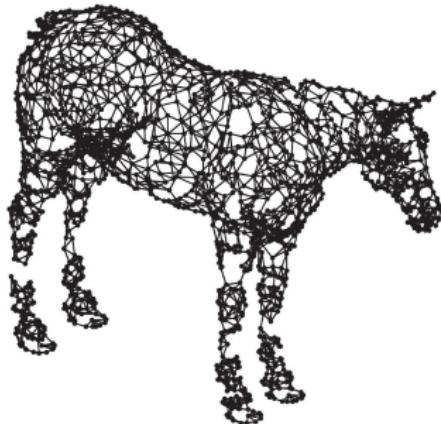
Vertices $\mathcal{V} = \{1, \dots, n\}$

Edges $\mathcal{E} \subseteq \mathcal{V} \times \mathcal{V}$

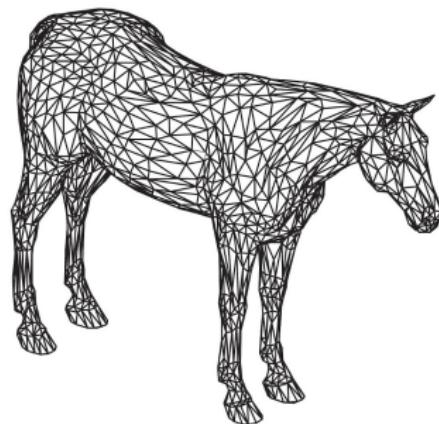
Edges $\mathcal{E} \subseteq \mathcal{V} \times \mathcal{V}$

Faces $\mathcal{F} = \{(i, j, k) \in \mathcal{V} \times \mathcal{V} \times \mathcal{V} : (i, j), (j, k), (k, i) \in \mathcal{E}\}$

Discrete manifolds



Nearest neighbor graph



Triangular mesh

Vertices $\mathcal{V} = \{1, \dots, n\}$

Vertices $\mathcal{V} = \{1, \dots, n\}$

Edges $\mathcal{E} \subseteq \mathcal{V} \times \mathcal{V}$

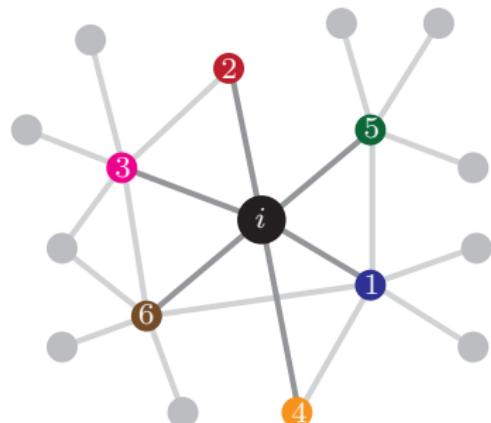
Edges $\mathcal{E} \subseteq \mathcal{V} \times \mathcal{V}$

Faces $\mathcal{F} = \{(i, j, k) \in \mathcal{V} \times \mathcal{V} \times \mathcal{V} : (i, j), (j, k), (k, i) \in \mathcal{E}\}$

Manifold mesh = each edge is shared by 2 faces + each vertex has 1 loop

Local ambiguity

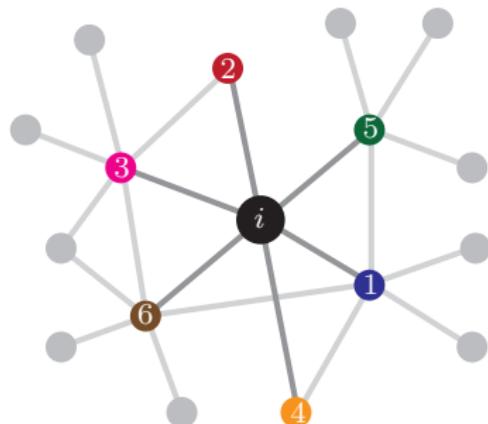
Unlike images, there is **no canonical ordering** of the domain points.



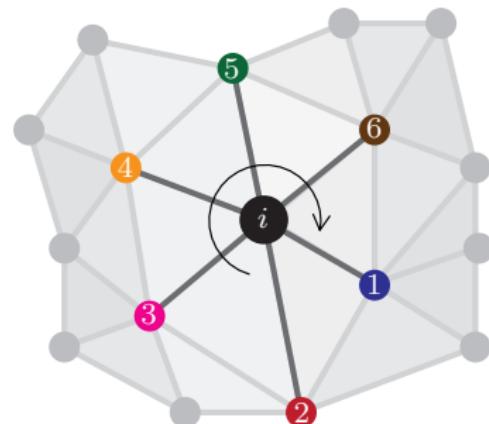
Graph (permutation)

Local ambiguity

Unlike images, there is **no canonical ordering** of the domain points.

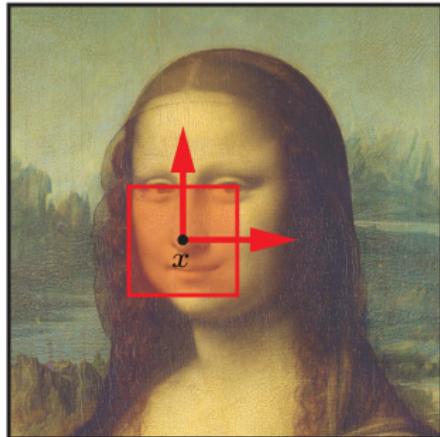


Graph (permutation)

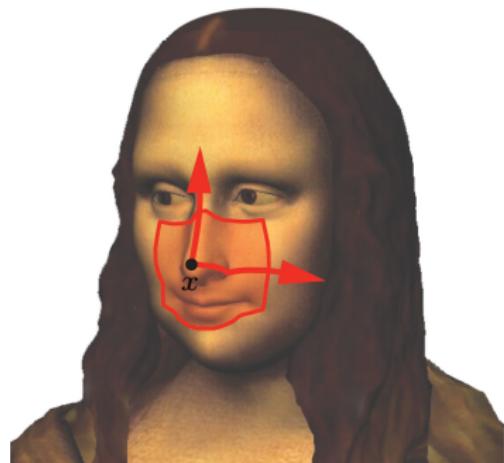


Mesh (rotation)

Non-Euclidean convolution?

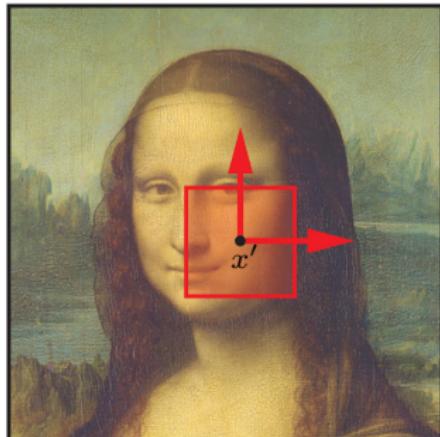


Euclidean

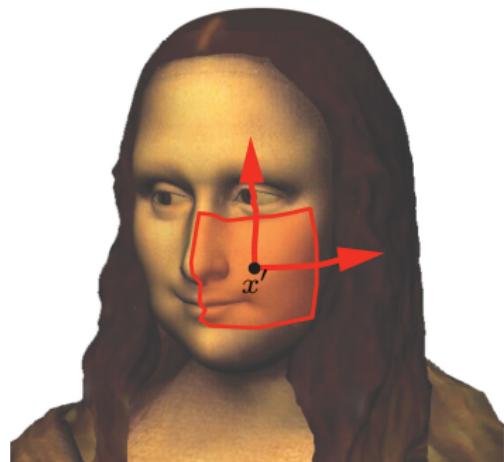


Non-Euclidean

Non-Euclidean convolution?

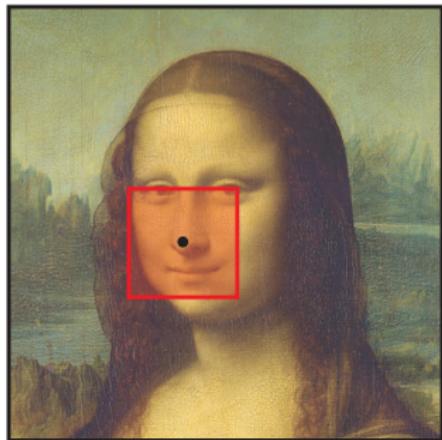


Euclidean

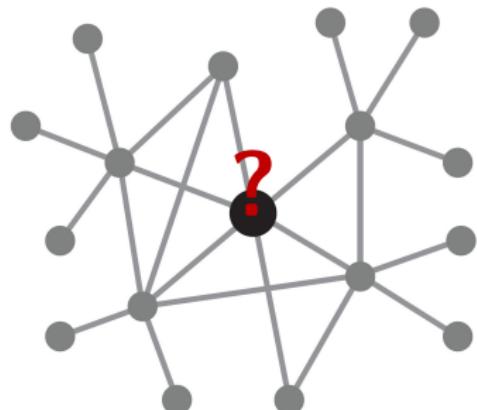


Non-Euclidean

Non-Euclidean convolution?



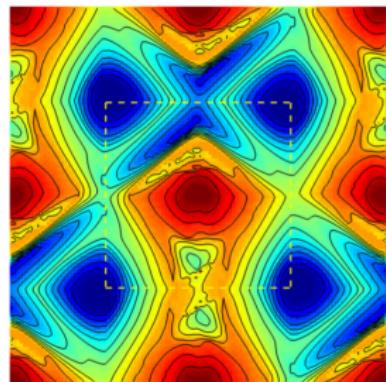
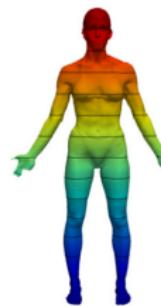
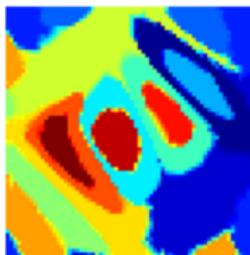
Image



Graph

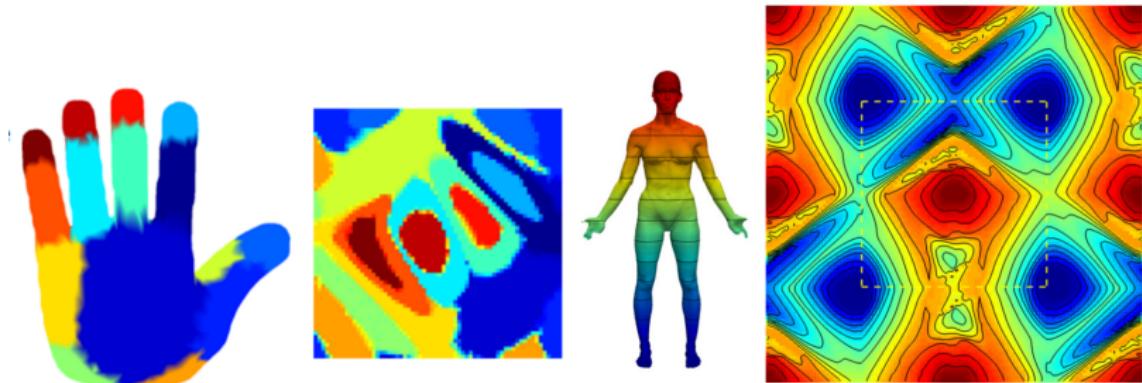
Global parametrization

Map the input mesh to some **parametric domain** (e.g. 2D plane) where operations can be defined more easily.



Global parametrization

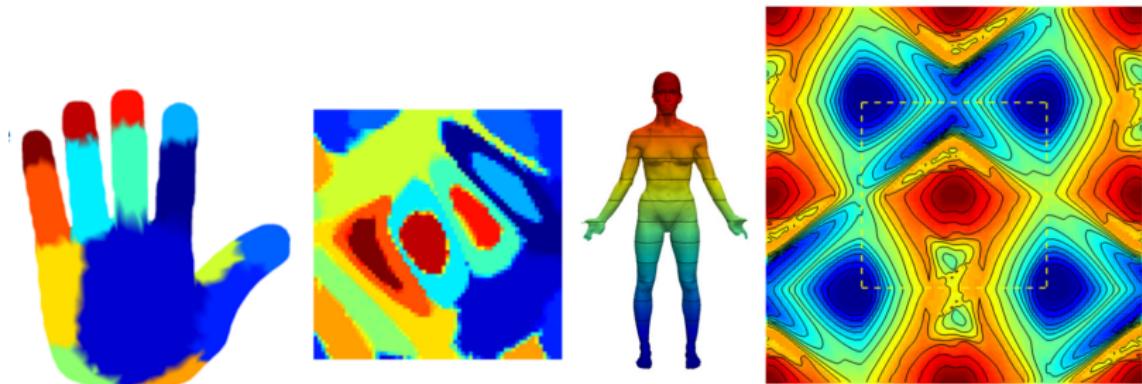
Map the input mesh to some **parametric domain** (e.g. 2D plane) where operations can be defined more easily.



- Can use **Euclidean** techniques in the embedding space

Global parametrization

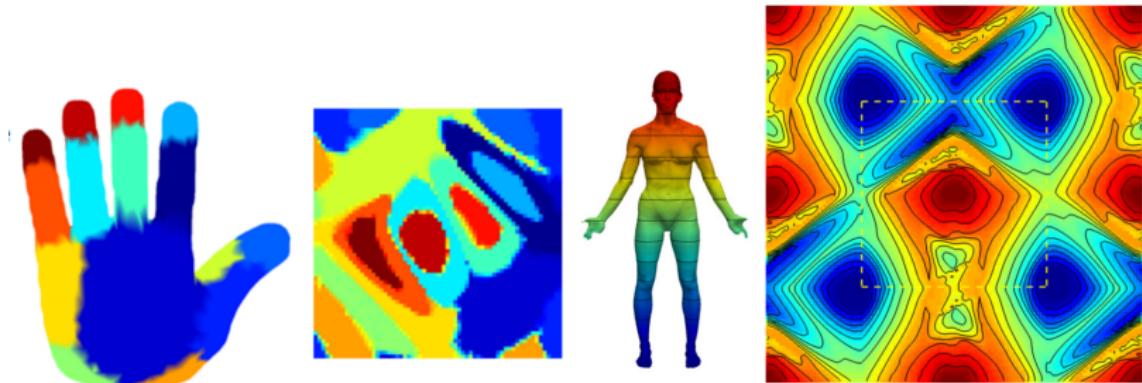
Map the input mesh to some **parametric domain** (e.g. 2D plane) where operations can be defined more easily.



- Can use **Euclidean** techniques in the embedding space
- Provides **invariance** to certain transformations

Global parametrization

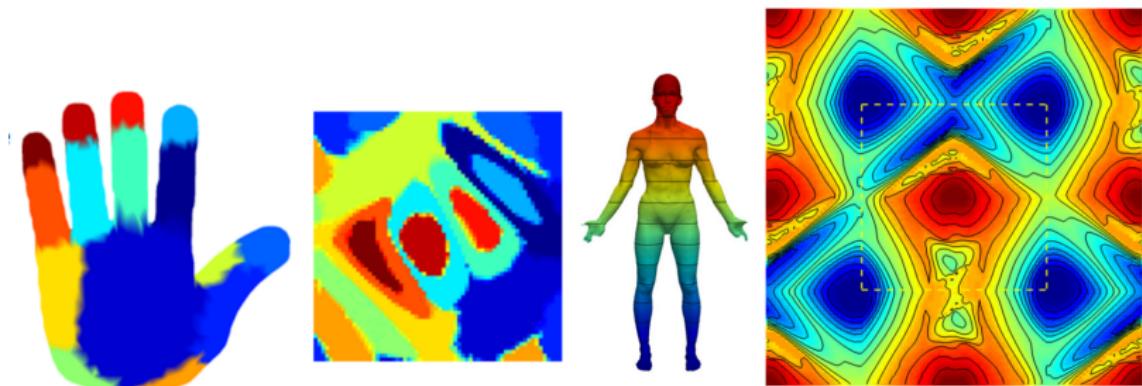
Map the input mesh to some **parametric domain** (e.g. 2D plane) where operations can be defined more easily.



- Can use **Euclidean** techniques in the embedding space
- Provides **invariance** to certain transformations
- Parametrization may be **non-unique**

Global parametrization

Map the input mesh to some **parametric domain** (e.g. 2D plane) where operations can be defined more easily.

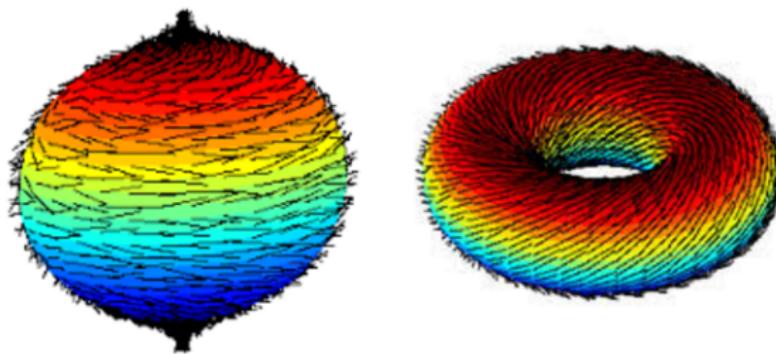


- Can use **Euclidean** techniques in the embedding space
- Provides **invariance** to certain transformations
- Parametrization may be **non-unique**
- The map can introduce **distortion**

Convolution on surfaces

Is **translation-invariant** convolution on surfaces possible?

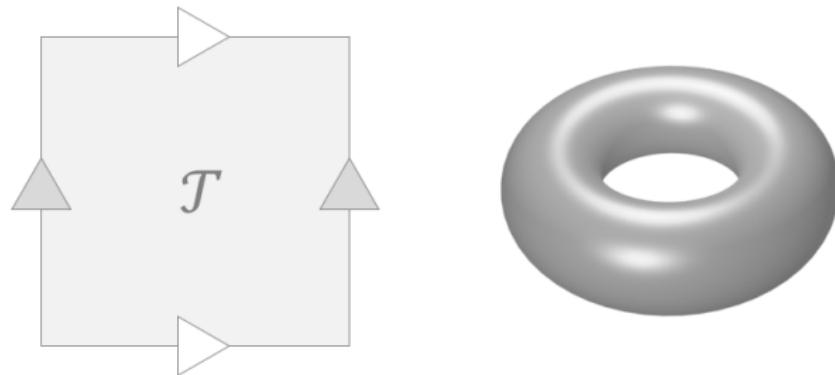
Not in general due to **singularities** in the translation field (Poincaré-Hopf or “hairy ball” theorem):



Convolution on surfaces

Is **translation-invariant** convolution on surfaces possible?

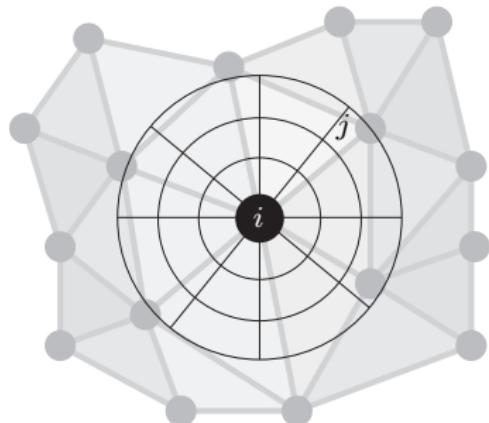
The **torus** is the only closed orientable surface admitting a translational group.



Convolution on surfaces

Spatial convolution on meshes

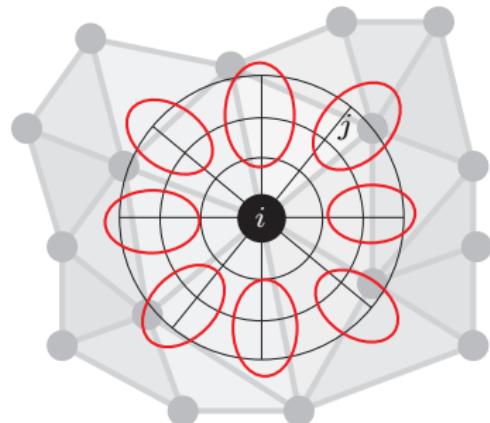
- Local system of coordinates \mathbf{u}_{ij} around i (e.g. geodesic polar)



Monti et al, "Geometric deep learning on graphs and manifolds using mixture model CNNs", CVPR 2016

Spatial convolution on meshes

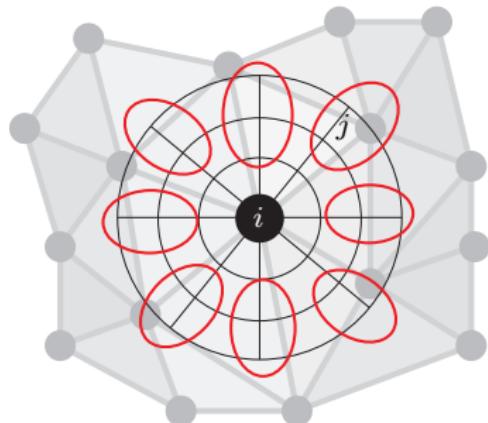
- Local system of coordinates \mathbf{u}_{ij} around i (e.g. geodesic polar)
- Local weights $w(\mathbf{u}_{ij})$, e.g. Gaussians with learnable μ, Σ :
$$w = \exp\left(-(\mathbf{u}_{ij} - \boldsymbol{\mu})^\top \boldsymbol{\Sigma}^{-1} (\mathbf{u}_{ij} - \boldsymbol{\mu})\right)$$



Monti et al, "Geometric deep learning on graphs and manifolds using mixture model CNNs", CVPR 2016

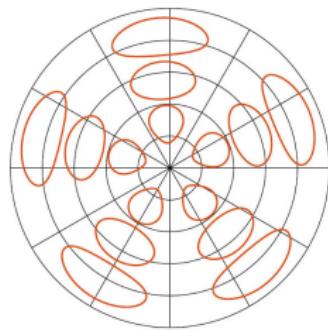
Spatial convolution on meshes

- Local system of coordinates \mathbf{u}_{ij} around i (e.g. geodesic polar)
- Local weights $w(\mathbf{u}_{ij})$, e.g. Gaussians with learnable μ, Σ :
$$w = \exp\left(-(\mathbf{u}_{ij} - \mu)^\top \Sigma^{-1} (\mathbf{u}_{ij} - \mu)\right)$$
- Spatial convolution of feature f with filter g :
 - Represent the input f as above
 $\Rightarrow \mathbf{f}$
 - Represent the learnable filter g as above $\Rightarrow \mathbf{g}$
 - Sum up the element-wise products $\Rightarrow \mathbf{f}^\top \mathbf{g}$

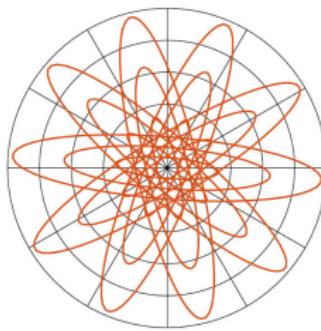


Monti et al, "Geometric deep learning on graphs and manifolds using mixture model CNNs", CVPR 2016

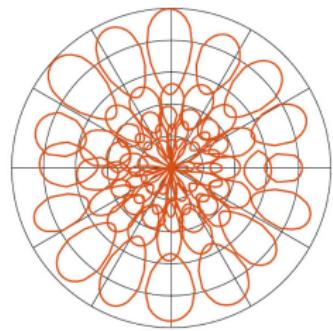
Local weighting kernels



GCNN



ACNN



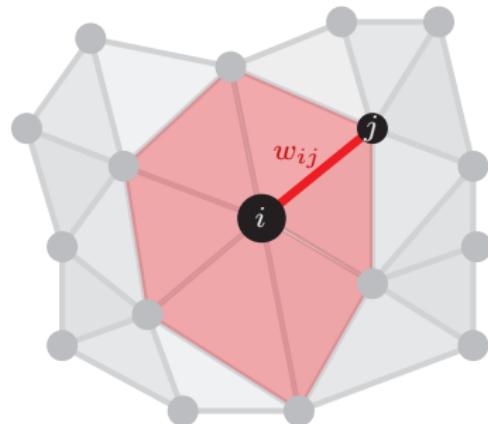
MoNet

Monti et al, "Geometric deep learning on graphs and manifolds using mixture model CNNs", CVPR 2016

Spectral convolution on meshes

- Laplacian operator Δ acting locally on the neighborhood of i :

$$(\Delta \mathbf{x})_i = \sum_j w_{ij} (\mathbf{x}_j - \mathbf{x}_i)$$

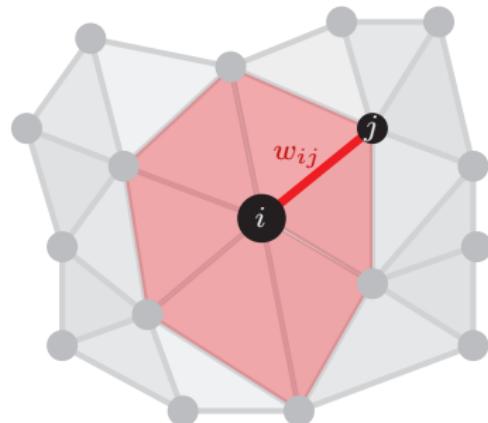


Spectral convolution on meshes

- Laplacian operator Δ acting locally on the neighborhood of i :

$$(\Delta \mathbf{x})_i = \sum_j w_{ij} (\mathbf{x}_j - \mathbf{x}_i)$$

= neighborhood avg – value at i



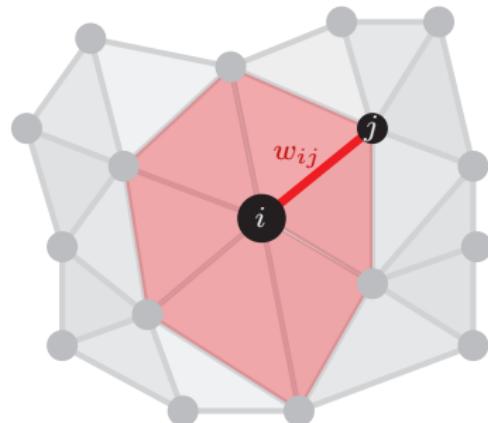
Spectral convolution on meshes

- Laplacian operator Δ acting locally on the neighborhood of i :

$$(\Delta \mathbf{x})_i = \sum_j w_{ij} (\mathbf{x}_j - \mathbf{x}_i)$$

- Eigenvectors of the Laplacian $\Delta = \Phi \Lambda \Phi^\top$ are a generalization of the Fourier transform:

$$\hat{\mathbf{x}} = \Phi^\top \mathbf{x}$$



Spectral convolution on meshes

- Laplacian operator Δ acting locally on the neighborhood of i :

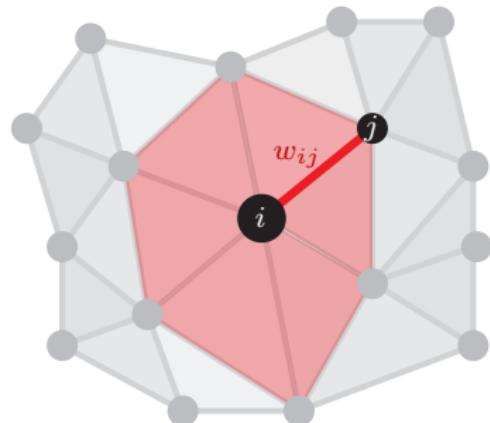
$$(\Delta \mathbf{x})_i = \sum_j w_{ij} (\mathbf{x}_j - \mathbf{x}_i)$$

- Eigenvectors of the Laplacian $\Delta = \Phi \Lambda \Phi^\top$ are a generalization of the Fourier transform:

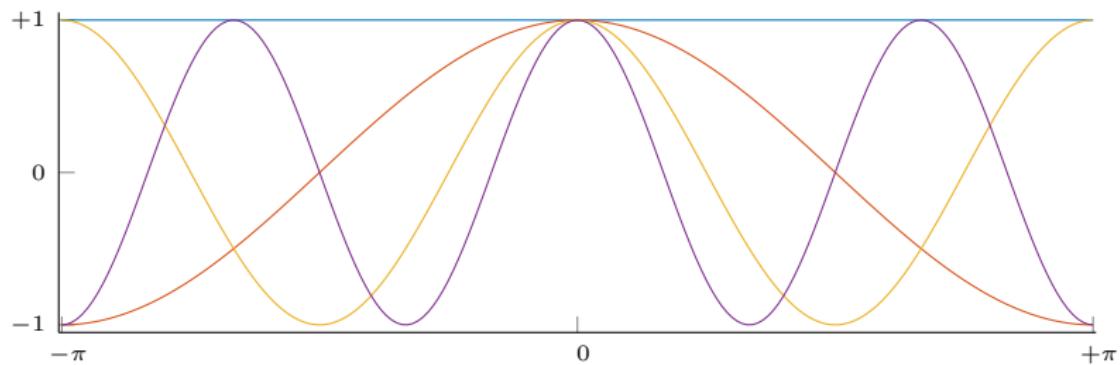
$$\hat{\mathbf{x}} = \Phi^\top \mathbf{x}$$

- Spectral convolution

$$\mathbf{x} * \mathbf{y} = \Phi \underbrace{\begin{pmatrix} \hat{y}_1 \\ \ddots \\ \hat{y}_n \end{pmatrix}}_{\hat{\mathbf{Y}}} \hat{\mathbf{x}}$$

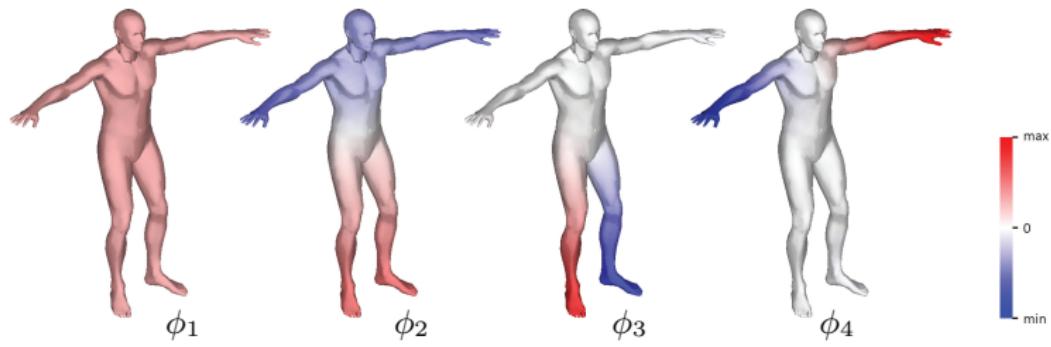


Laplacian eigenfunctions: Euclidean



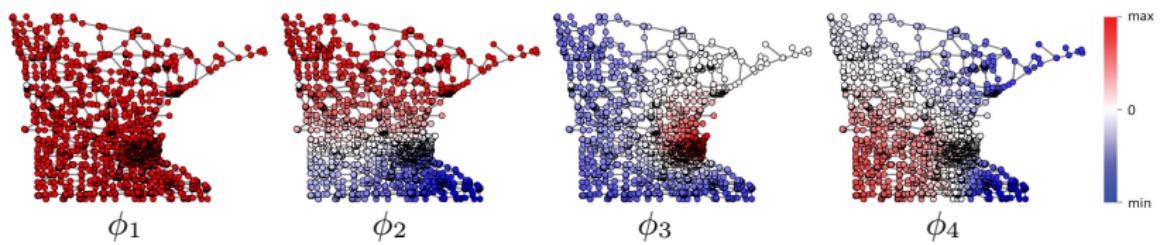
First eigenfunctions of 1D Euclidean Laplacian = standard Fourier basis

Laplacian eigenfunctions: manifold



First eigenfunctions of a manifold Laplacian

Laplacian eigenfunctions: graph

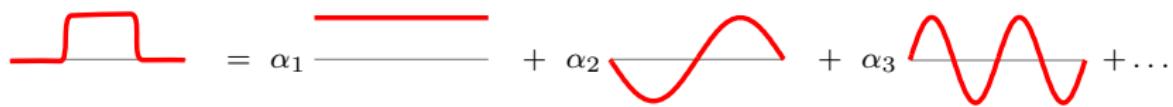


First eigenfunctions of a graph Laplacian

Fourier analysis: Euclidean space

A function $f : [-\pi, \pi] \rightarrow \mathbb{R}$ can be written as Fourier series

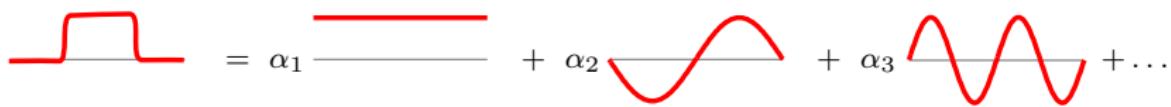
$$f(x) = \sum_{k \geq 0} \frac{1}{2\pi} \int_{-\pi}^{\pi} f(x') e^{-ikx'} dx' e^{ikx}$$



Fourier analysis: Euclidean space

A function $f : [-\pi, \pi] \rightarrow \mathbb{R}$ can be written as Fourier series

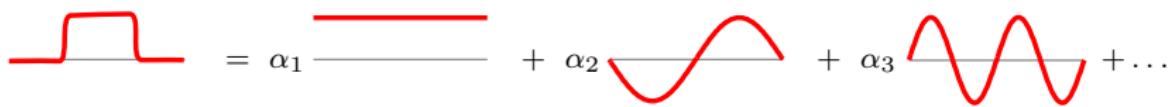
$$f(x) = \sum_{k \geq 0} \underbrace{\frac{1}{2\pi} \int_{-\pi}^{\pi} f(x') e^{-ikx'} dx'}_{\hat{f}_k = \langle f, e^{ikx} \rangle_{L^2([-\pi, \pi])}} e^{ikx}$$



Fourier analysis: Euclidean space

A function $f : [-\pi, \pi] \rightarrow \mathbb{R}$ can be written as Fourier series

$$f(x) = \sum_{k \geq 0} \underbrace{\frac{1}{2\pi} \int_{-\pi}^{\pi} f(x') e^{-ikx'} dx'}_{\hat{f}_k = \langle f, e^{ikx} \rangle_{L^2([-\pi, \pi])}} e^{ikx}$$

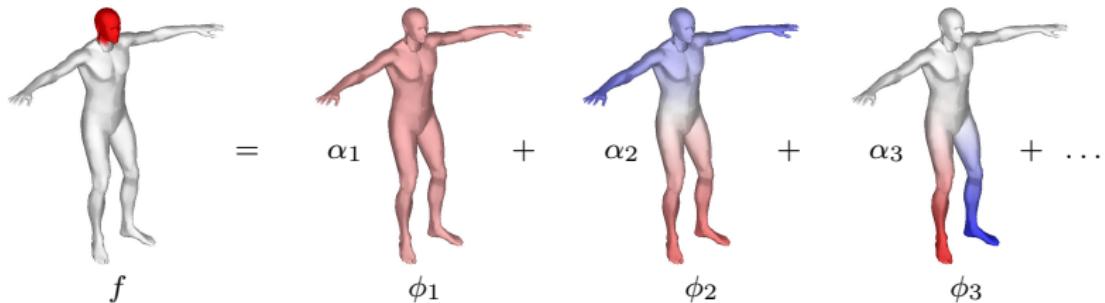


Fourier basis = Laplacian eigenfunctions: $-\frac{d^2}{dx^2} e^{ikx} = k^2 e^{ikx}$

Fourier analysis: non-Euclidean space

A function $f : \mathcal{X} \rightarrow \mathbb{R}$ can be written as Fourier series

$$f(x) = \sum_{k \geq 1} \underbrace{\int_{\mathcal{X}} f(x') \phi_k(x') dx'}_{\hat{f}_k = \langle f, \phi_k \rangle_{L^2(\mathcal{X})}} \phi_k(x)$$



Fourier basis = Laplacian eigenfunctions: $\Delta \phi_k(x) = \lambda_k \phi_k(x)$

Convolution theorem

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

Convolution theorem

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

Convolution theorem: Fourier transform **diagonalizes** the convolution operator \Rightarrow convolution can be computed in the Fourier domain as:

$$\widehat{(f \star g)} = \hat{f} \cdot \hat{g}$$

Convolution theorem

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

Convolution of two vectors $\mathbf{f} = (f_1, \dots, f_n)^\top$ and $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\mathbf{f} \star \mathbf{g} = \begin{bmatrix} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{bmatrix} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix}$$

Convolution theorem

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

Convolution of two vectors $\mathbf{f} = (f_1, \dots, f_n)^\top$ and $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\mathbf{f} \star \mathbf{g} = \underbrace{\begin{bmatrix} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{bmatrix}}_{\text{circulant matrix}} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix}$$

Convolution theorem

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

Convolution of two vectors $\mathbf{f} = (f_1, \dots, f_n)^\top$ and $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\mathbf{f} \star \mathbf{g} = \underbrace{\begin{bmatrix} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{bmatrix}}_{\text{diagonalized by Fourier basis}} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix}$$

Convolution theorem

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

Convolution of two vectors $\mathbf{f} = (f_1, \dots, f_n)^\top$ and $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\mathbf{f} \star \mathbf{g} = \begin{bmatrix} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{bmatrix} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix}$$

$$= \Phi \begin{bmatrix} \hat{g}_1 & & & \\ & \ddots & & \\ & & \ddots & \\ & & & \hat{g}_n \end{bmatrix} \Phi^\top \mathbf{f}$$

Convolution theorem

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

Convolution of two vectors $\mathbf{f} = (f_1, \dots, f_n)^\top$ and $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\mathbf{f} \star \mathbf{g} = \begin{bmatrix} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{bmatrix} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix}$$

$$= \Phi \begin{bmatrix} \hat{g}_1 & & & \\ & \ddots & & \\ & & \hat{g}_n & \end{bmatrix} \begin{bmatrix} \hat{f}_1 \\ \vdots \\ \hat{f}_n \end{bmatrix}$$

Convolution theorem

Given two functions $f, g : [-\pi, \pi] \rightarrow \mathbb{R}$ their **convolution** is a function

$$(f \star g)(x) = \int_{-\pi}^{\pi} f(x')g(x - x')dx'$$

Convolution of two vectors $\mathbf{f} = (f_1, \dots, f_n)^\top$ and $\mathbf{g} = (g_1, \dots, g_n)^\top$

$$\mathbf{f} \star \mathbf{g} = \begin{bmatrix} g_1 & g_2 & \dots & \dots & g_n \\ g_n & g_1 & g_2 & \dots & g_{n-1} \\ \vdots & \vdots & \ddots & \ddots & \vdots \\ g_3 & g_4 & \dots & g_1 & g_2 \\ g_2 & g_3 & \dots & \dots & g_1 \end{bmatrix} \begin{bmatrix} f_1 \\ \vdots \\ f_n \end{bmatrix}$$

$$= \Phi \begin{bmatrix} \hat{f}_1 \cdot \hat{g}_1 \\ \vdots \\ \hat{f}_n \cdot \hat{g}_n \end{bmatrix}$$

Spectral convolution

Generalized convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})} \phi_k$$

Spectral convolution

Generalized convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \underbrace{\langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})}}_{\text{product in the Fourier domain}} \phi_k$$

Spectral convolution

Generalized convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \underbrace{\langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})}}_{\text{product in the Fourier domain}} \underbrace{\phi_k}_{\text{inverse Fourier transform}}$$

Spectral convolution

Generalized convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})} \phi_k$$

In matrix-vector notation

$$\mathbf{f} \star \mathbf{g} = \Phi (\Phi^\top \mathbf{g}) \circ (\Phi^\top \mathbf{f})$$

Spectral convolution

Generalized convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})} \phi_k$$

In matrix-vector notation

$$\mathbf{f} \star \mathbf{g} = \Phi \operatorname{diag}(\hat{g}_1, \dots, \hat{g}_n) \Phi^\top \mathbf{f}$$

Spectral convolution

Generalized convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})} \phi_k$$

In matrix-vector notation

$$\mathbf{f} \star \mathbf{g} = \underbrace{\Phi \operatorname{diag}(\hat{g}_1, \dots, \hat{g}_n) \Phi^\top}_{\mathbf{G}} \mathbf{f}$$

Spectral convolution

Generalized convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

$$f \star g = \sum_{k \geq 1} \langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})} \phi_k$$

In matrix-vector notation

$$\mathbf{f} \star \mathbf{g} = \underbrace{\Phi \operatorname{diag}(\hat{g}_1, \dots, \hat{g}_n) \Phi^\top}_{\mathbf{G}} \mathbf{f}$$

- Not shift-invariant! (\mathbf{G} has no circulant structure)

Spectral convolution

Generalized convolution of $f, g \in L^2(\mathcal{X})$ can be defined by analogy

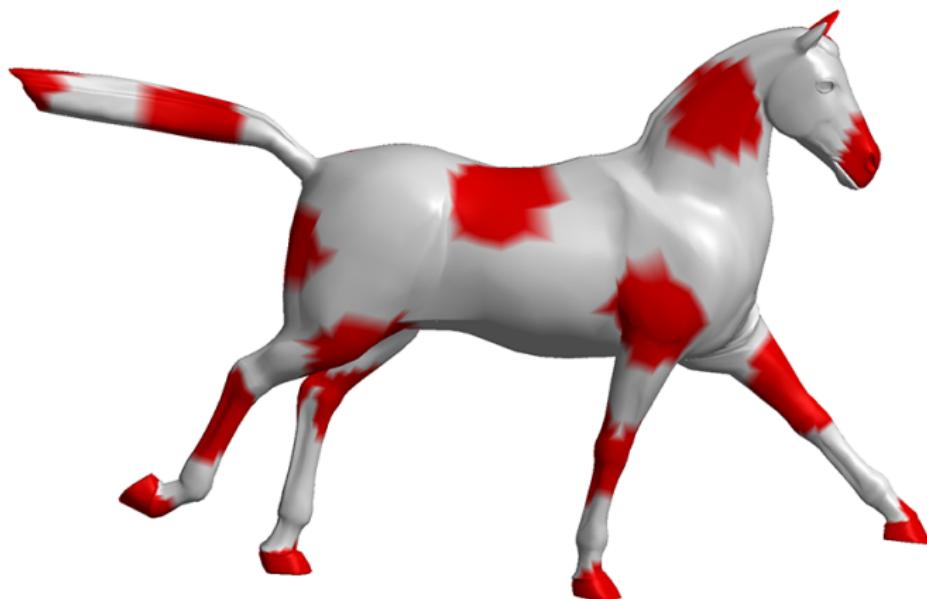
$$f \star g = \sum_{k \geq 1} \langle f, \phi_k \rangle_{L^2(\mathcal{X})} \langle g, \phi_k \rangle_{L^2(\mathcal{X})} \phi_k$$

In matrix-vector notation

$$\mathbf{f} \star \mathbf{g} = \underbrace{\Phi \operatorname{diag}(\hat{g}_1, \dots, \hat{g}_n) \Phi^\top}_{\mathbf{G}} \mathbf{f}$$

- Not shift-invariant! (\mathbf{G} has no circulant structure)
- Filter coefficients depend on basis ϕ_1, \dots, ϕ_n

Basis dependence



Function x

Basis dependence



'Edge detecting' spectral filter $\hat{\Phi} \hat{Y} \Phi^\top x$

Basis dependence

Same spectral filter, different basis $\Psi \hat{Y} \Psi^\top x$

Spectral convolution on meshes

- Laplacian operator Δ acting locally on the neighborhood of i :

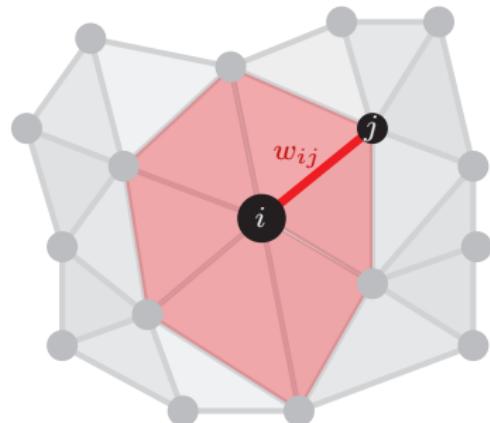
$$(\Delta \mathbf{x})_i = \sum_j w_{ij} (\mathbf{x}_j - \mathbf{x}_i)$$

- Eigenvectors of the Laplacian $\Delta = \Phi \Lambda \Phi^\top$ are a generalization of the Fourier transform:

$$\hat{\mathbf{x}} = \Phi^\top \mathbf{x}$$

- Spectral convolution:

$$\mathbf{x} \star \mathbf{y} = \Phi \underbrace{\begin{pmatrix} \hat{y}_1 \\ \ddots \\ \hat{y}_n \end{pmatrix}}_{\hat{\mathbf{Y}}} \hat{\mathbf{x}}$$



Spectral convolution on meshes

- Laplacian operator Δ acting locally on the neighborhood of i :

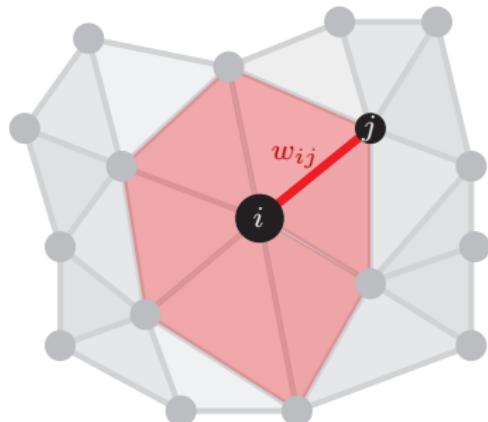
$$(\Delta \mathbf{x})_i = \sum_j w_{ij} (\mathbf{x}_j - \mathbf{x}_i)$$

- Eigenvectors of the Laplacian $\Delta = \Phi \Lambda \Phi^\top$ are a generalization of the Fourier transform:

$$\hat{\mathbf{x}} = \Phi^\top \mathbf{x}$$

- Spectral convolution defined as a filter applied on the Laplacian:

$$\mathbf{X}' = \Phi \tau(\Lambda) \Phi^\top \mathbf{X}$$



Locality and smoothness

In the Euclidean setting (by Parseval's identity), the following holds:

$$\int_{-\infty}^{+\infty} |x|^{2k} |f(x)|^2 dx = \int_{-\infty}^{+\infty} \left| \frac{\partial^k \hat{f}(\omega)}{\partial \omega^k} \right|^2 d\omega$$

Localization in space = smoothness in frequency domain

Locality and smoothness

In the Euclidean setting (by Parseval's identity), the following holds:

$$\int_{-\infty}^{+\infty} |x|^{2k} |f(x)|^2 dx = \int_{-\infty}^{+\infty} \left| \frac{\partial^k \hat{f}(\omega)}{\partial \omega^k} \right|^2 d\omega$$

Localization in space = smoothness in frequency domain

Parametrize the filter using a smooth spectral transfer function $\tau(\lambda)$.

Locality and smoothness

In the Euclidean setting (by Parseval's identity), the following holds:

$$\int_{-\infty}^{+\infty} |x|^{2k} |f(x)|^2 dx = \int_{-\infty}^{+\infty} \left| \frac{\partial^k \hat{f}(\omega)}{\partial \omega^k} \right|^2 d\omega$$

Localization in space = smoothness in frequency domain

Parametrize the filter using a smooth spectral transfer function $\tau(\lambda)$.

Application of the filter

$$\tau(\Delta)\mathbf{f} = \Phi \tau(\Lambda) \Phi^\top \mathbf{f}$$

Locality and smoothness

In the Euclidean setting (by Parseval's identity), the following holds:

$$\int_{-\infty}^{+\infty} |x|^{2k} |f(x)|^2 dx = \int_{-\infty}^{+\infty} \left| \frac{\partial^k \hat{f}(\omega)}{\partial \omega^k} \right|^2 d\omega$$

Localization in space = smoothness in frequency domain

Parametrize the filter using a smooth spectral transfer function $\tau(\lambda)$.

Application of the filter

$$\tau(\Delta)\mathbf{f} = \Phi \begin{pmatrix} \tau(\lambda_1) & & \\ & \ddots & \\ & & \tau(\lambda_n) \end{pmatrix} \Phi^\top \mathbf{f}$$

Locality and smoothness

In the Euclidean setting (by Parseval's identity), the following holds:

$$\int_{-\infty}^{+\infty} |x|^{2k} |f(x)|^2 dx = \int_{-\infty}^{+\infty} \left| \frac{\partial^k \hat{f}(\omega)}{\partial \omega^k} \right|^2 d\omega$$

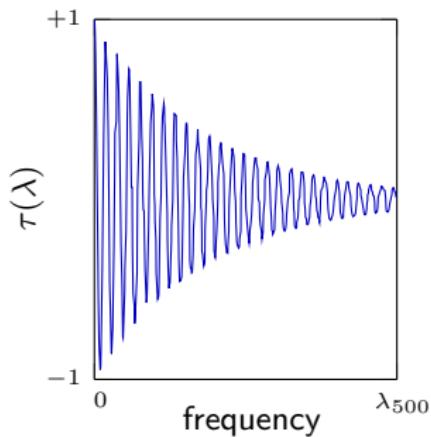
Localization in space = smoothness in frequency domain

Parametrize the filter using a smooth spectral transfer function $\tau(\lambda)$.

Application of the parametric filter with learnable parameters α

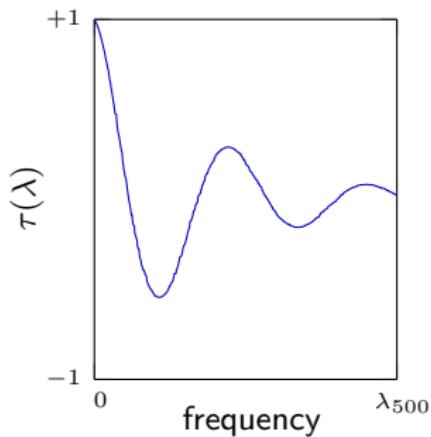
$$\tau_\alpha(\Delta)\mathbf{f} = \Phi \begin{pmatrix} \tau_\alpha(\lambda_1) & & \\ & \ddots & \\ & & \tau_\alpha(\lambda_n) \end{pmatrix} \Phi^\top \mathbf{f}$$

Locality and smoothness



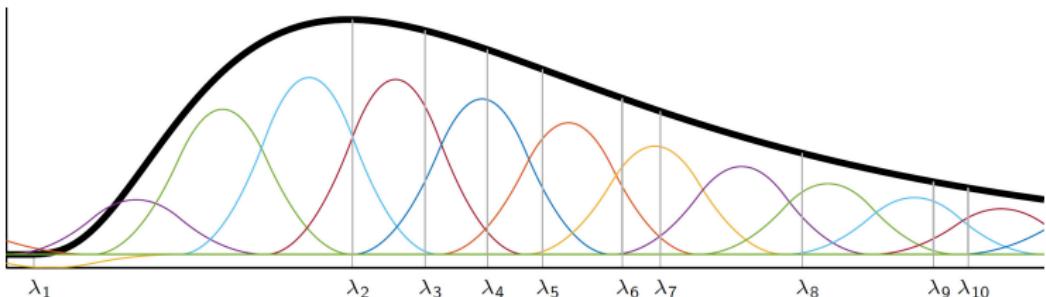
Non-smooth spectral filter (delocalized in space)

Locality and smoothness



Smooth spectral filter (localized in space)

Spectral graph CNN with smooth spectral filters

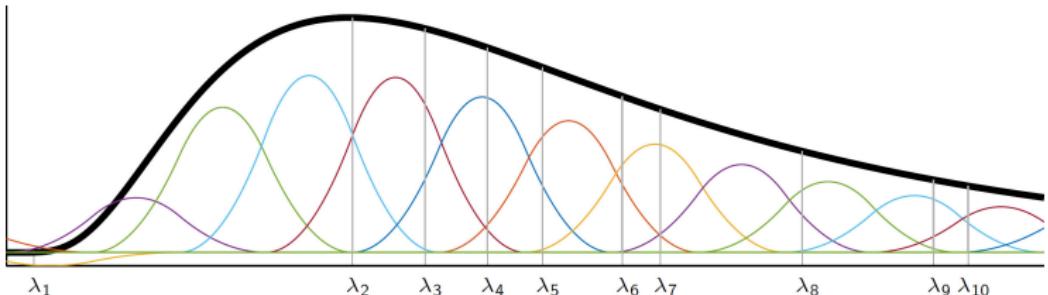


Consider a linear combination of smooth kernel functions
 $\beta_1(\lambda), \dots, \beta_r(\lambda)$:

$$\tau_{\alpha}(\lambda) = \sum_{j=1}^r \alpha_j \beta_j(\lambda)$$

where $\alpha = (\alpha_1, \dots, \alpha_r)^{\top}$ is the vector of filter parameters.

Spectral graph CNN with smooth spectral filters

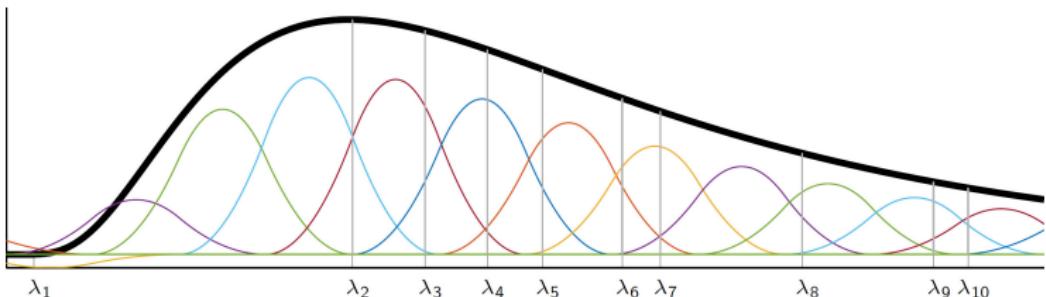


Consider a linear combination of smooth kernel functions
 $\beta_1(\lambda), \dots, \beta_r(\lambda)$:

$$\tau_{\alpha}(\lambda_k) = \sum_{j=1}^r \alpha_j \beta_j(\lambda_k)$$

where $\alpha = (\alpha_1, \dots, \alpha_r)^{\top}$ is the vector of filter parameters.

Spectral graph CNN with smooth spectral filters

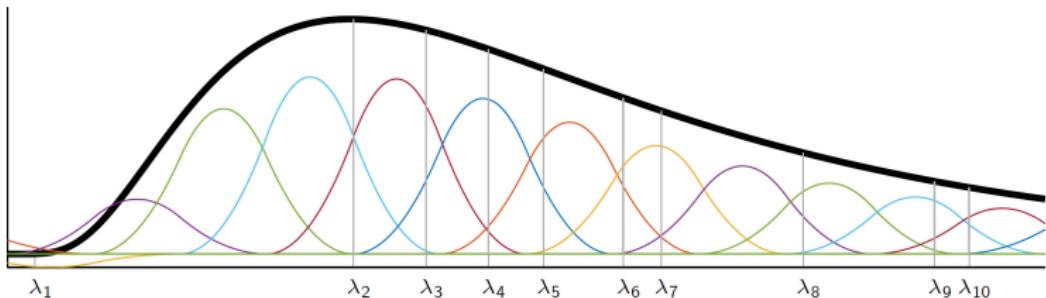


Consider a linear combination of smooth kernel functions
 $\beta_1(\lambda), \dots, \beta_r(\lambda)$:

$$\tau_{\boldsymbol{\alpha}}(\lambda_k) = \sum_{j=1}^r \alpha_j \beta_j(\lambda_k) = (\mathbf{B}\boldsymbol{\alpha})_k$$

where $\boldsymbol{\alpha} = (\alpha_1, \dots, \alpha_r)^\top$ is the vector of filter parameters.

Spectral graph CNN with smooth spectral filters

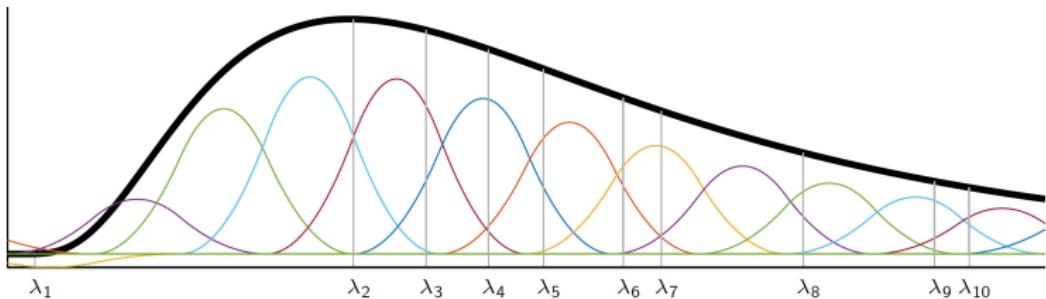


Consider a **linear combination of smooth kernel functions**
 $\beta_1(\lambda), \dots, \beta_r(\lambda)$:

$$\mathbf{W} = \text{Diag}(\mathbf{B}\boldsymbol{\alpha})$$

where $\boldsymbol{\alpha} = (\alpha_1, \dots, \alpha_r)^\top$ is the vector of filter parameters.

Spectral graph CNN with smooth spectral filters



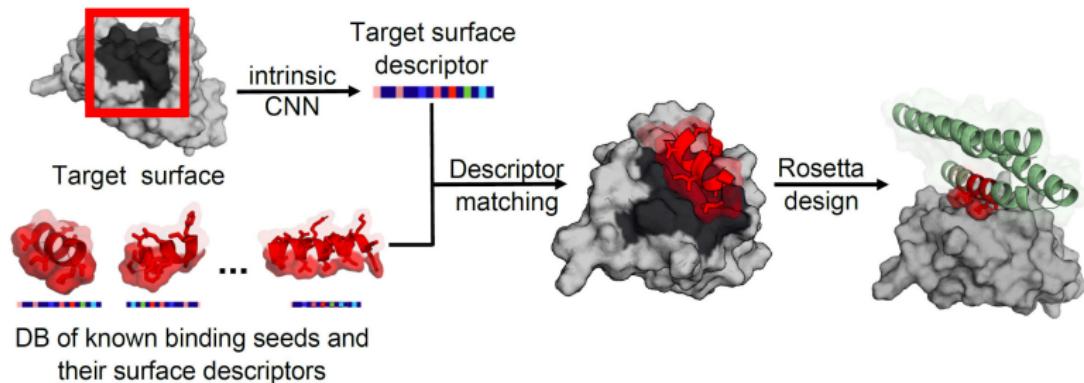
Consider a **linear combination of smooth kernel functions**
 $\beta_1(\lambda), \dots, \beta_r(\lambda)$:

$$\mathbf{W} = \text{Diag}(\mathbf{B}\boldsymbol{\alpha})$$

where $\boldsymbol{\alpha} = (\alpha_1, \dots, \alpha_r)^\top$ is the vector of filter parameters.

$\mathcal{O}(1)$ parameters per layer.

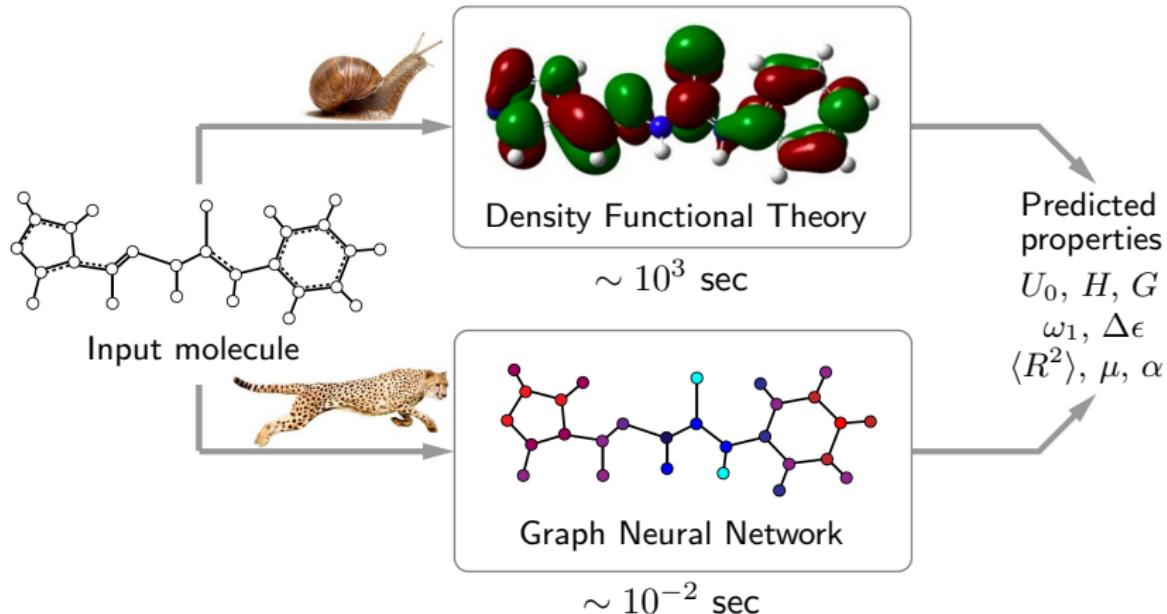
Application: Protein-Protein Interaction



Designing protein binder for the PD-L1 protein target

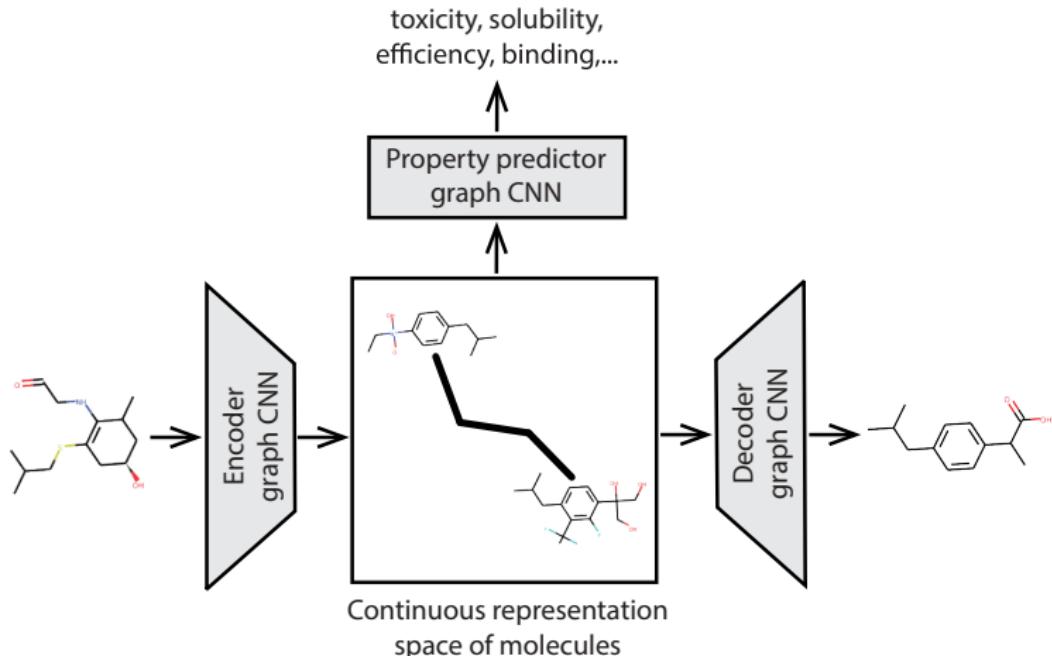
Gainza et al, "Deciphering interaction fingerprints from protein molecular surfaces using geometric deep learning", Nature Methods 2020

Molecule property prediction

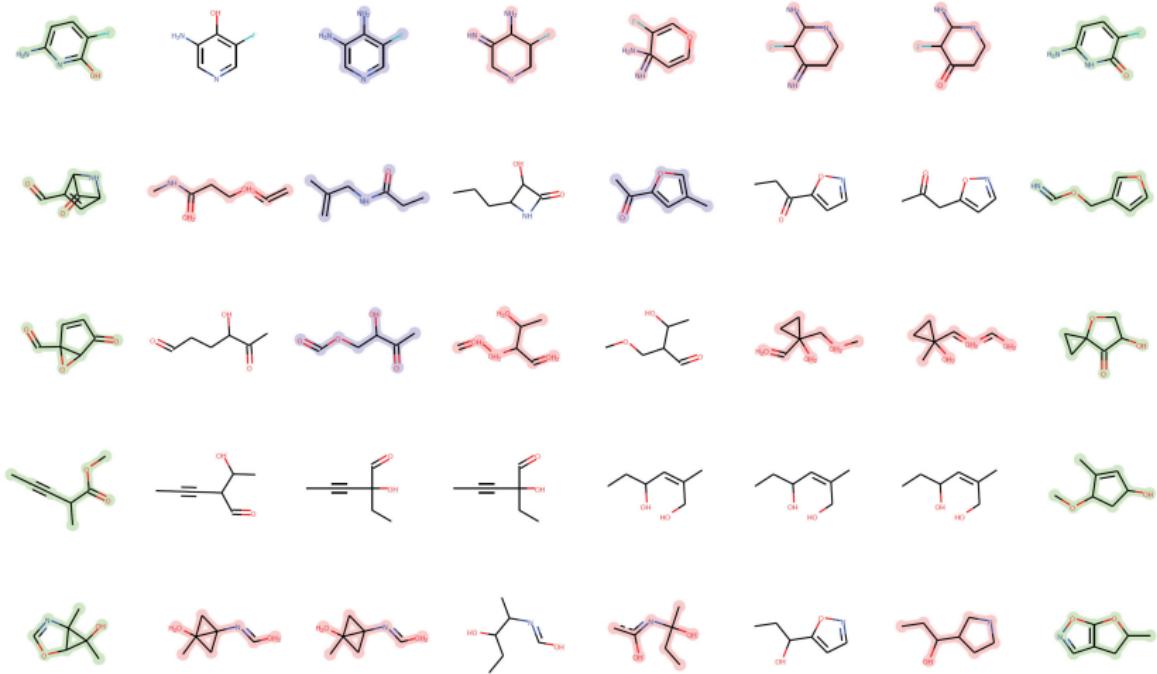


Duvenaud et al, "Convolutional Networks on Graphs for Learning Molecular Fingerprints", NIPS 2015; Gomez-Bombarelli et al, "Automatic chemical design using a data-driven continuous representation of molecules", ACS Cent. Sci. 2018

Generative models



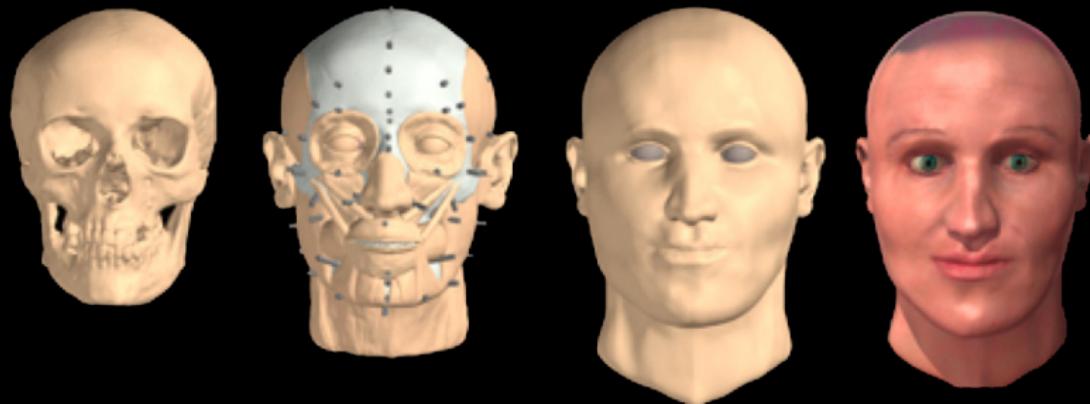
Molecule generation



Molecules generated with a graph VAE

Simonovsky and Komodakis, "Graphvae: Towards generation of small graphs using variational autoencoders", 2017; De Cao and Kipf, "MolGAN: An implicit generative model for small molecular graphs", 2018

Face from DNA



Claes et al, "Facial recognition from DNA using face-to-DNA classifiers", Nature Communications 2019

Suggested reading

Bronstein et al, 2016

“Geometric deep learning: going beyond Euclidean data”

<https://arxiv.org/abs/1611.08097>